Recommendations_with_IBM

February 10, 2023

1 Recommendations with IBM

In this notebook, you will be putting your recommendation skills to use on real data from the IBM Watson Studio platform.

You may either submit your notebook through the workspace here, or you may work from your local machine and submit through the next page. Either way assure that your code passes the project RUBRIC. **Please save regularly.**

By following the table of contents, you will build out a number of different methods for making recommendations that can be used for different situations.

1.1 Table of Contents

At the end of the notebook, you will find directions for how to submit your work. Let's get started by importing the necessary libraries and reading in the data.

```
In [1]: import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        import seaborn as sns
        import project_tests as t
        import pickle

        %matplotlib inline

        df = pd.read_csv('data/user-item-interactions.csv')
        df_content = pd.read_csv('data/articles_community.csv')
        del df['Unnamed: 0']
        del df_content['Unnamed: 0']

        # Show df to get an idea of the data
        df.head()

Out[1]:    article_id                                              title  \
        0      1430.0  using pixiedust for fast, flexible, and easier...
        1      1314.0       healthcare python streaming application demo
        2      1429.0          use deep learning for image classification
        3      1338.0            ml optimization using cognitive assistant
```

1

```
        4       1276.0          deploy your python model as a restful api

                                                         email
        0  ef5f11f77ba020cd36e1105a00ab868bbdbf7fe7
        1  083cbdfa93c8444beaa4c5f5e0f5f9198e4f9e0b
        2  b96a4f2e92d8572034b1e9b28f9ac673765cd074
        3  06485706b34a5c9bf2a0ecdac41daf7e7654ceb7
        4  f01220c46fc92c6e6b161b1849de11faacd7ccb2
```

In [2]: *# Show df_content to get an idea of the data*
        df_content.head()

Out[2]:   doc_body \
        0 Skip navigation Sign in SearchLoading...\r\n\r...
        1 No Free Hunch Navigation * kaggle.com\r\n\r\n ...
        2   * Login\r\n * Sign Up\r\n\r\n * Learning Pat...
        3 DATALAYER: HIGH THROUGHPUT, LOW LATENCY AT SCA...
        4 Skip navigation Sign in SearchLoading...\r\n\r...

                                              doc_description \
        0 Detect bad readings in real time using Python ...
        1 See the forest, see the trees. Here lies the c...
        2 Heres this weeks news in Data Science and Bi...
        3 Learn how distributed DBs solve the problem of...
        4 This video demonstrates the power of IBM DataS...

                                   doc_full_name doc_status article_id
        0 Detect Malfunctioning IoT Sensors with Streami...   Live   0
        1 Communicating data science: A guide to present...   Live   1
        2   This Week in Data Science (April 18, 2017)   Live   2
        3 DataLayer Conference: Boost the performance of...   Live   3
        4   Analyze NY Restaurant data using Spark in DSX   Live   4

1.1.1 Part I : Exploratory Data Analysis

Use the dictionary and cells below to provide some insight into the descriptive statistics of the
data.

   1. What is the distribution of how many articles a user interacts with in the dataset? Provide
a visual and descriptive statistics to assist with giving a look at the number of times each user
interacts with an article.

In [3]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45993 entries, 0 to 45992
Data columns (total 3 columns):
article_id    45993 non-null float64
title         45993 non-null object
email         45976 non-null object
```

```
dtypes: float64(1), object(2)
memory usage: 1.1+ MB
```

In [4]: df.shape

Out[4]: (45993, 3)

In [5]: df.groupby(["email","article_id"]).max().unstack().head(3) #email-article interaction (m

```
Out[5]:                                               title                                   \
        article_id                              0.0    2.0    4.0    8.0    9.0
        email
        0000b6387a0366322d7fbfc6434af145adf7fed1   NaN    NaN    NaN    NaN    NaN
        001055fc0bb67f71e8fa17002342b256a30254cd   NaN    NaN    NaN    NaN    NaN
        00148e4911c7e04eeff8def7bbbdaf1c59c2c621   NaN    NaN    NaN    NaN    NaN


                                                                                 \
        article_id                              12.0   14.0   15.0   16.0   18.0
        email
        0000b6387a0366322d7fbfc6434af145adf7fed1   NaN    NaN    NaN    NaN    NaN
        001055fc0bb67f71e8fa17002342b256a30254cd   NaN    NaN    NaN    NaN    NaN
        00148e4911c7e04eeff8def7bbbdaf1c59c2c621   NaN    NaN    NaN    NaN    NaN


                                                 ...                            \
        article_id                              ...   1434.0 1435.0 1436.0 1437.0
        email                                   ...
        0000b6387a0366322d7fbfc6434af145adf7fed1 ...    NaN    NaN    NaN    NaN
        001055fc0bb67f71e8fa17002342b256a30254cd ...    NaN    NaN    NaN    NaN
        00148e4911c7e04eeff8def7bbbdaf1c59c2c621 ...    NaN    NaN    NaN    NaN


                                                                                 \
        article_id                              1439.0 1440.0 1441.0 1442.0 1443.0
        email
        0000b6387a0366322d7fbfc6434af145adf7fed1   NaN    NaN    NaN    NaN    NaN
        001055fc0bb67f71e8fa17002342b256a30254cd   NaN    NaN    NaN    NaN    NaN
        00148e4911c7e04eeff8def7bbbdaf1c59c2c621   NaN    NaN    NaN    NaN    NaN


        article_id                              1444.0
        email
        0000b6387a0366322d7fbfc6434af145adf7fed1   NaN
        001055fc0bb67f71e8fa17002342b256a30254cd   NaN
        00148e4911c7e04eeff8def7bbbdaf1c59c2c621   NaN

        [3 rows x 714 columns]
```

Explore the Data

```
In [6]: user_article_i = df.groupby("email").count()
        user_article_i.head(6)

Out[6]:                                              article_id  title
        email
        0000b6387a0366322d7fbfc6434af145adf7fed1            13     13
        001055fc0bb67f71e8fa17002342b256a30254cd             4      4
        00148e4911c7e04eeff8def7bbbdaf1c59c2c621             3      3
        001a852ecbd6cc12ab77a785efa137b2646505fe             6      6
        001fc95b90da5c3cb12c501d201a915e4f093290             2      2
        0042719415c4fca7d30bd2d4e9d17c5fc570de13             2      2

In [7]: user_article_i.describe()

Out[7]:           article_id         title
        count   5148.000000   5148.000000
        mean       8.930847      8.930847
        std       16.802267     16.802267
        min        1.000000      1.000000
        25%        1.000000      1.000000
        50%        3.000000      3.000000
        75%        9.000000      9.000000
        max      364.000000    364.000000

In [8]: user_article_i[user_article_i["article_id"]==364]

Out[8]:                                              article_id  title
        email
        2b6c0f514c2f2b04ad3c4583407dccd0810469ee           364    364

In [9]: sns.distplot(user_article_i["article_id"],bins = 100, kde = False) # Distribution of the
        plt.title('Distribution of User Article Interactions')
        plt.xlabel('# of Interactions');
```

Distribution of User Article Interactions

In [10]: # Fill in the median and maximum number of user_article interactios below

```python
median_val = df.groupby("email").count()["article_id"].median()
max_views_by_user = df.groupby("email").count()["article_id"].max()

print(median_val)
print("\n")
print(max_views_by_user)
```

3.0

364

2. Explore and remove duplicate articles from the **df_content** dataframe.

In [11]: df_content.shape # shape of the dataset

Out[11]: (1056, 5)

In [12]: df_content[df_content.duplicated("article_id")]

Out[12]:   doc_body \
         365   Follow Sign in / Sign up Home About Insight Da...

```
692  Homepage Follow Sign in / Sign up Homepage * H...
761  Homepage Follow Sign in Get started Homepage *...
970  This video shows you how to construct queries ...
971  Homepage Follow Sign in Get started * Home\r\n...

                                    doc_description  \
365  During the seven-week Insight Data Engineering...
692  One of the earliest documented catalogs was co...
761  Todays world of data science leverages data f...
970  This video shows you how to construct queries ...
971  If you are like most data scientists, you are ...

                                    doc_full_name doc_status  article_id
365                    Graph-based machine learning       Live          50
692  How smart catalogs can turn the big data flood...       Live         221
761  Using Apache Spark as a parallel processing fr...       Live         398
970                            Use the Primary Index       Live         577
971  Self-service data preparation with IBM Data Re...       Live         232
```

In [13]: *# Remove any rows that have the same article_id - only keep the first*
         df_content.drop_duplicates(subset="article_id",inplace=True,keep="first")

In [14]: *# After droping dublicates*
         df_content.shape

Out[14]: (1051, 5)

3. Use the cells below to find:
a. The number of unique articles that have an interaction with a user.
b. The number of unique articles in the dataset (whether they have any interactions or not). **c.** The number of unique users in the dataset. (excluding null values) **d.** The number of user-article interactions in the dataset.

In [15]: unique_articles = df["article_id"].nunique() *# The number of unique articles that have*
         total_articles = df_content["article_id"].nunique() *# The number of unique articles on*
         unique_users = df["email"].nunique() *# The number of unique users*
         user_article_interactions = df.shape[0]

         print(unique_articles)
         print(total_articles)
         print(unique_users)
         print(user_article_interactions)

714
1051
5148
45993

4. Use the cells below to find the most viewed **article_id**, as well as how often it was viewed. After talking to the company leaders, the email_mapper function was deemed a reasonable way to map users to ids. There were a small number of null values, and it was found that all of these null values likely belonged to a single user (which is how they are stored using the function below).

```
In [16]: most_viewed_article_id = str(df["article_id"].value_counts().index[0])  # Most viewed a
         max_views = df.groupby("article_id").count().max()["email"] # Max Views

In [17]: ## No need to change the code here - this will be helpful for later parts of the notebo
         # Run this cell to map the user email to a user_id column and remove the email column

         def email_mapper():
             coded_dict = dict()
             cter = 1
             email_encoded = []

             for val in df['email']:
                 if val not in coded_dict:
                     coded_dict[val] = cter
                     cter+=1

                 email_encoded.append(coded_dict[val])
             return email_encoded

         email_encoded = email_mapper()
         del df['email']
         df['user_id'] = email_encoded

         # show header
         df.head()
```

```
Out[17]:    article_id                                            title  user_id
         0     1430.0  using pixiedust for fast, flexible, and easier...        1
         1     1314.0         healthcare python streaming application demo        2
         2     1429.0            use deep learning for image classification        3
         3     1338.0            ml optimization using cognitive assistant        4
         4     1276.0            deploy your python model as a restful api        5
```

```
In [18]: ## If you stored all your results in the variable names above,
         ## you shouldn't need to change anything in this cell

         sol_1_dict = {
             '`50% of individuals have _____ or fewer interactions.`': median_val,
             '`The total number of user-article interactions in the dataset is _____.`': user_a
             '`The maximum number of user-article interactions by any 1 user is _____.`': max_v
             '`The most viewed article in the dataset was viewed _____ times.`': max_views,
             '`The article_id of the most viewed article is _____.`': most_viewed_article_id,
             '`The number of unique articles that have at least 1 rating _____.`': unique_artic
             '`The number of unique users in the dataset is _____`': unique_users,
```

```
            '`The number of unique articles on the IBM platform`': total_articles
        }

        # Test your dictionary against the solution
        t.sol_1_test(sol_1_dict)
```

It looks like you have everything right here! Nice job!

1.1.2 Part II: Rank-Based Recommendations

Unlike in the earlier lessons, we don't actually have ratings for whether a user liked an article or not. We only know that a user has interacted with an article. In these cases, the popularity of an article can really only be based on how often an article was interacted with.

1. Fill in the function below to return the **n** top articles ordered with most interactions as the top. Test your function using the tests below.

```
In [19]: def get_top_articles(n, df=df):
             '''
             INPUT:
             n - (int) the number of top articles to return
             df - (pandas dataframe) df as defined at the top of the notebook

             OUTPUT:
             top_articles - (list) A list of the top 'n' article titles

             '''
             top_articles = df["title"].value_counts()[:n].index.tolist()
             top_articles = [str(i) for i in top_articles]

             return top_articles # Return the top article titles from df (not df_content)

         def get_top_article_ids(n, df=df):
             '''
             INPUT:
             n - (int) the number of top articles to return
             df - (pandas dataframe) df as defined at the top of the notebook

             OUTPUT:
             top_articles - (list) A list of the top 'n' article titles

             '''
             top_articles = df["article_id"].value_counts().index.tolist()[:n]
             top_articles = [str(i) for i in top_articles]

             return top_articles # Return the top article ids

In [20]: print(get_top_articles(10))
         print(get_top_article_ids(10))
```

8

```
['use deep learning for image classification', 'insights from new york car accident reports', 'v
['1429.0', '1330.0', '1431.0', '1427.0', '1364.0', '1314.0', '1293.0', '1170.0', '1162.0', '1304
```

```
In [21]: # Test your function by returning the top 5, 10, and 20 articles
         top_5 = get_top_articles(5)
         top_10 = get_top_articles(10)
         top_20 = get_top_articles(20)

         # Test each of your three lists from above
         t.sol_2_test(get_top_articles)
```

```
Your top_5 looks like the solution list! Nice job.
Your top_10 looks like the solution list! Nice job.
Your top_20 looks like the solution list! Nice job.
```

1.1.3 Part III: User-User Based Collaborative Filtering

1. Use the function below to reformat the **df** dataframe to be shaped with users as the rows and articles as the columns.

- Each **user** should only appear in each **row** once.

- Each **article** should only show up in one **column**.

- **If a user has interacted with an article, then place a 1 where the user-row meets for that article-column**. It does not matter how many times a user has interacted with the article, all entries where a user has interacted with an article should be a 1.

- **If a user has not interacted with an item, then place a zero where the user-row meets for that article-column**.

Use the tests to make sure the basic structure of your matrix matches what is expected by the solution.

```
In [22]: # create the user-article matrix with 1's and 0's

         def create_user_item_matrix(df):
             '''
             INPUT:
             df - pandas dataframe with article_id, title, user_id columns

             OUTPUT:
             user_item - user item matrix

             Description:
             Return a matrix with user ids as rows and article ids on the columns with 1 values
             an article and a 0 otherwise
             '''
```

```python
                # Fill in the function here
                user_item = df.groupby(["user_id","article_id"]).count().unstack()
                user_item.fillna(0, inplace = True)
                user_item = user_item.applymap(lambda x: 1 if x>0 else x )

                return user_item["title"] # return the user_item matrix

         user_item = create_user_item_matrix(df)

In [23]: ## Tests: You should just need to run this cell.  Don't change the code.
         assert user_item.shape[0] == 5149, "Oops!  The number of users in the user-article matr
         assert user_item.shape[1] == 714, "Oops!  The number of articles in the user-article ma
         assert user_item.sum(axis=1)[1] == 36, "Oops!  The number of articles seen by user 1 do
         print("You have passed our quick tests!  Please proceed!")

You have passed our quick tests!  Please proceed!
```

2. Complete the function below which should take a user_id and provide an ordered list of the most similar users to that user (from most similar to least similar). The returned result should not contain the provided user_id, as we know that each user is similar to him/herself. Because the results for each user here are binary, it (perhaps) makes sense to compute similarity as the dot product of two users.

Use the tests to test your function.

```python
In [24]: def find_similar_users(user_id, user_item=user_item):
             '''
             INPUT:
             user_id - (int) a user_id
             user_item - (pandas dataframe) matrix of users by articles:
                         1's when a user has interacted with an article, 0 otherwise

             OUTPUT:
             similar_users - (list) an ordered list where the closest users (largest dot product
                             are listed first

             Description:
             Computes the similarity of every pair of users based on the dot product
             Returns an ordered

             '''
             # compute similarity of each user to the provided user
             simi_matrix = user_item.dot(np.transpose(user_item))

             # sort by similarity
             simi_matrix = simi_matrix.loc[user_id].sort_values(ascending = False)

             # create list of just the ids
             most_similar_users = simi_matrix.index.tolist()
```

```
                # remove the own user's id
                most_similar_users.remove(user_id)
                return most_similar_users # return a list of the users in order from most to least
```

In [25]: `# Do a spot check of your function`
```
         print("The 10 most similar users to user 1 are: {}".format(find_similar_users(1)[:10]))
         print("The 5 most similar users to user 3933 are: {}".format(find_similar_users(3933)[:
         print("The 3 most similar users to user 46 are: {}".format(find_similar_users(46)[:3]))
```

```
The 10 most similar users to user 1 are: [3933, 23, 3782, 203, 4459, 131, 3870, 46, 4201, 5041]
The 5 most similar users to user 3933 are: [1, 23, 3782, 4459, 203]
The 3 most similar users to user 46 are: [4201, 23, 3782]
```

In [26]: `df_content.head() # for reference`

Out[26]:
```
                                                      doc_body  \
         0   Skip navigation Sign in SearchLoading...\r\n\r...
         1   No Free Hunch Navigation * kaggle.com\r\n\r\n ...
         2    * Login\r\n * Sign Up\r\n\r\n * Learning Pat...
         3   DATALAYER: HIGH THROUGHPUT, LOW LATENCY AT SCA...
         4   Skip navigation Sign in SearchLoading...\r\n\r...

                                                doc_description  \
         0   Detect bad readings in real time using Python ...
         1   See the forest, see the trees. Here lies the c...
         2   Heres this weeks news in Data Science and Bi...
         3   Learn how distributed DBs solve the problem of...
         4   This video demonstrates the power of IBM DataS...

                                       doc_full_name doc_status   article_id
         0   Detect Malfunctioning IoT Sensors with Streami...       Live            0
         1   Communicating data science: A guide to present...       Live            1
         2           This Week in Data Science (April 18, 2017)       Live            2
         3   DataLayer Conference: Boost the performance of...       Live            3
         4        Analyze NY Restaurant data using Spark in DSX       Live            4
```

3. Now that you have a function that provides the most similar users to each user, you will want to use these users to find articles you can recommend. Complete the functions below to return the articles you would recommend to each user.

In [27]: `def get_article_names(article_ids, df=df):`
```
             '''
             INPUT:
             article_ids - (list) a list of article ids
             df - (pandas dataframe) df as defined at the top of the notebook
```

11

```python
    OUTPUT:
    article_names - (list) a list of article names associated with the list of article
                    (this is identified by the title column)
    '''
    article_names = list()
    df = df.set_index("article_id")

    for idd in article_ids:
        article_names.append(df.loc[float(idd)].max()["title"])

    return article_names

def get_user_articles(user_id, user_item=user_item):
    '''
    INPUT:
    user_id - (int) a user id
    user_item - (pandas dataframe) matrix of users by articles:
                1's when a user has interacted with an article, 0 otherwise

    OUTPUT:
    article_ids - (list) a list of the article ids seen by the user
    article_names - (list) a list of article names associated with the list of article

    Description:
    Provides a list of the article_ids and article titles that have been seen by a user
    '''
    article_ids = user_item.loc[user_id][user_item.loc[user_id].values == 1].index.asty

    article_names = []

    for idd in article_ids:
        article_names.append(df[df['article_id']==float(idd)].max()['title']) # need to

    return article_ids, article_names


def user_user_recs(user_id, m = 10):
    '''
    INPUT:
    user_id - (int) a user id
    m - (int) the number of recommendations you want for the user

    OUTPUT:
    recs - (list) a list of recommendations for the user

    Description:
    Loops through the users based on closeness to the input user_id
    For each user - finds articles the user hasn't seen before and provides them as rec
```

```
        Does this until m recommendations are found

        Notes:
        Users who are the same closeness are chosen arbitrarily as the 'next' user

        For the user where the number of recommended articles starts below m
        and ends exceeding m, the last items are chosen arbitrarily

        '''

        neighbor_users = find_similar_users(user_id) # closest neighbor to our user_id
        user_articles = get_user_articles(user_id)[0] #seen by our user

        recs = np.array([])


        for user in neighbor_users:

            neighbor_articles_seen = get_user_articles(user)[0] # movies seen by others lik
            recs1 = np.setdiff1d(neighbor_articles_seen, user_articles, assume_unique=True)
            recs = np.concatenate([recs1, recs], axis = 0) #concanate arrays
            recs = np.unique(recs) # find unique items in array

            if len(recs) > m-1:
                break

        recs = recs[:m]
        recs.tolist()

        return recs # return your recommendations for this user_id
```

In [28]: `# Check Results`
`get_article_names(user_user_recs(1, 10)) # Return 10 recommendations for user 1`

Out[28]: ['recommender systems: approaches & algorithms',
         '1448   i ranked every intro to data science course on...\nName: title, dtype: object
         'data tidying in data science experience',
         'a tensorflow regression model to predict house values',
         '520   using notebooks with pixiedust for fast, flexi...\nName: title, dtype: object'
         'airbnb data for analytics: mallorca reviews',
         'airbnb data for analytics: vancouver listings',
         'analyze facebook data using ibm watson and watson studio',
         'analyze accident reports on amazon emr spark',
         'analyze energy consumption in buildings']

In [29]: `# Test your functions here - No need to change this code - just run this cell`
`assert set(get_article_names(['1024.0', '1176.0', '1305.0', '1314.0', '1422.0', '1427.0`
`assert set(get_article_names(['1320.0', '232.0', '844.0'])) == set(['housing (2015): un`

13

```
        assert set(get_user_articles(20)[0]) == set(['1320.0', '232.0', '844.0'])
        assert set(get_user_articles(20)[1]) == set(['housing (2015): united states demographic
        assert set(get_user_articles(2)[0]) == set(['1024.0', '1176.0', '1305.0', '1314.0', '14
        assert set(get_user_articles(2)[1]) == set(['using deep learning to reconstruct high-re
        print("If this is all you see, you passed all of our tests!  Nice job!")

If this is all you see, you passed all of our tests!  Nice job!
```

4. Now we are going to improve the consistency of the **user_user_recs** function from above.

- Instead of arbitrarily choosing when we obtain users who are all the same closeness to a given user - choose the users that have the most total article interactions before choosing those with fewer article interactions.

- Instead of arbitrarily choosing articles from the user where the number of recommended articles starts below m and ends exceeding m, choose articles with the articles with the most total interactions before choosing those with fewer total interactions. This ranking should be what would be obtained from the **top_articles** function you wrote earlier.

```
In [30]: def get_top_sorted_users(user_id, df=df, user_item=user_item):
             '''
             INPUT:
             user_id - (int)
             df - (pandas dataframe) df as defined at the top of the notebook
             user_item - (pandas dataframe) matrix of users by articles:
                     1's when a user has interacted with an article, 0 otherwise


             OUTPUT:
             neighbors_df - (pandas dataframe) a dataframe with:
                             neighbor_id - is a neighbor user_id
                             similarity - measure of the similarity of each user to the provided
                             num_interactions - the number of articles viewed by the user - if a

             Other Details - sort the neighbors_df by the similarity and then by number of inter
                             highest of each is higher in the dataframe

             '''
             # dot product of user_item with user_item transpose which gives the similarities bt
             user_dot_pro = user_item.dot(np.transpose(user_item))

             # dataframe with neighbor_id and similarity columns
             neighbors_df = user_dot_pro.loc[user_id].rename_axis("neighbor_id").reset_index(nam

             # dataframe with the user-article interactions
             interaction_df = df["user_id"].value_counts().rename_axis("neighbor_id").reset_inde

             # Merged dataframe of neighbors_df and interaction_df
```

14

```python
        neighbors_df = pd.merge(neighbors_df,interaction_df,on="neighbor_id",how = "outer")

        # sort the dataframe with similarity first and then number of interactions
        neighbors_df = neighbors_df.sort_values(by=['similarity', 'num_interactions'], asce

        # Remove the row with the input user
        neighbors_df = neighbors_df[neighbors_df["neighbor_id"] != user_id]

        return neighbors_df # Return the dataframe specified in the doc_string


def user_user_recs_part2(user_id, m=10):
    '''
    INPUT:
    user_id - (int) a user id
    m - (int) the number of recommendations you want for the user

    OUTPUT:
    recs - (list) a list of recommendations for the user by article id
    rec_names - (list) a list of recommendations for the user by article title

    Description:
    Loops through the users based on closeness to the input user_id
    For each user - finds articles the user hasn't seen before and provides them as rec
    Does this until m recommendations are found

    Notes:
    * Choose the users that have the most total article interactions
    before choosing those with fewer article interactions.

    * Choose articles with the articles with the most total interactions
    before choosing those with fewer total interactions.

    '''
    recs = np.array([])

    user_articles_ids_seen, user_articles_names_seen = get_user_articles(user_id, user_
    closest_neighs = get_top_sorted_users(user_id, df, user_item).neighbor_id.tolist()

    for neighs in closest_neighs:

        neigh_articles_ids_seen, neigh_articles_names_seen = get_user_articles(neighs,
        new_recs = np.setdiff1d(neigh_articles_ids_seen, user_articles_ids_seen, assume
        recs = np.unique(np.concatenate([new_recs, recs], axis = 0)) # concate arrays a

        if len(recs) > m-1:
            break
```

```
            recs = recs[:m]
            recs = recs.tolist() # convert to a list

            rec_names = get_article_names(recs, df=df)

            return recs, rec_names

            return recs, rec_names
```

In [31]: # Quick spot check - don't change this code - just use it to test your functions
         rec_ids, rec_names = user_user_recs_part2(20, 10)
         print("The top 10 recommendations for user 20 are the following article ids:")
         print(rec_ids)
         print()
         print("The top 10 recommendations for user 20 are the following article names:")
         print(rec_names)

The top 10 recommendations for user 20 are the following article ids:
['1024.0', '1085.0', '109.0', '1150.0', '1151.0', '1152.0', '1153.0', '1154.0', '1157.0', '1160.

The top 10 recommendations for user 20 are the following article names:
['using deep learning to reconstruct high-resolution audio', 'airbnb data for analytics: chicago

5. Use your functions from above to correctly fill in the solutions to the dictionary below. Then test your dictionary against the solution. Provide the code you need to answer each following the comments below.

In [32]: ### Tests with a dictionary of results

         user1_most_sim = get_top_sorted_users(1)["neighbor_id"].iloc[0]
         user131_10th_sim = get_top_sorted_users(131)["neighbor_id"].iloc[9]

In [33]: ## Dictionary Test Here
         sol_5_dict = {
             'The user that is most similar to user 1.': user1_most_sim,
             'The user that is the 10th most similar to user 131': user131_10th_sim,
         }

         t.sol_5_test(sol_5_dict)

This all looks good! Nice job!

6. If we were given a new user, which of the above functions would you be able to use to make recommendations? Explain. Can you think of a better way we might make recommendations? Use the cell below to explain a better method for new users.

   The function that is the best amongst the functions created is the knowledge based function (get_top_articles) as it gives the top articles interacted with. It is the best one for a new user as

we do not have enough information regarding the new user, to which the new user has interacted with. There can be a better way to recommend if we have enough information regarding the user and we can use the combination of knowledge, collaboration type recommendatio

Provide your response here.

7. Using your existing functions, provide the top 10 recommended articles you would provide for the a new user below. You can test your function against our thoughts to make sure we are all on the same page with how we might make a recommendation.

```
In [34]: new_user = '0.0'

         # What would your recommendations be for this new user '0.0'?  As a new user, they have
         # Provide a list of the top 10 article ids you would give to
         new_user_recs = get_top_article_ids(10, df=df)

In [35]: assert set(new_user_recs) == set(['1314.0','1429.0','1293.0','1427.0','1162.0','1364.0'

         print("That's right!  Nice job!")

That's right!  Nice job!
```

1.1.4   Part IV: Content Based Recommendations (EXTRA - NOT REQUIRED)

Another method we might use to make recommendations is to perform a ranking of the highest ranked articles associated with some term. You might consider content to be the **doc_body**, **doc_description**, or **doc_full_name**. There isn't one way to create a content based recommendation, especially considering that each of these columns hold content related information.

1. Use the function body below to create a content based recommender. Since there isn't one right answer for this recommendation tactic, no test functions are provided. Feel free to change the function inputs if you decide you want to try a method that requires more input values. The input values are currently set with one idea in mind that you may use to make content based recommendations. One additional idea is that you might want to choose the most popular recommendations that meet your 'content criteria', but again, there is a lot of flexibility in how you might make these recommendations.

1.1.5   This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.

```
In [36]: def make_content_recs():
             '''
             INPUT:

             OUTPUT:

             '''
             pass
```

2. Now that you have put together your content-based recommendation system, use the cell below to write a summary explaining how your content based recommender works. Do you see

any possible improvements that could be made to your function? Is there anything novel about your content based recommender?

1.1.6 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.

Write an explanation of your content based recommendation system here.

3. Use your content-recommendation system to make recommendations for the below scenarios based on the comments. Again no tests are provided here, because there isn't one right answer that could be used to find these content based recommendations.

1.1.7 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.

```
In [37]: # make recommendations for a brand new user


         # make a recommendations for a user who only has interacted with article id '1427.0'
```

1.1.8 Part V: Matrix Factorization

In this part of the notebook, you will build use matrix factorization to make article recommendations to the users on the IBM Watson Studio platform.

1. You should have already created a **user_item** matrix above in **question 1** of **Part III** above. This first question here will just require that you run the cells to get things set up for the rest of **Part V** of the notebook.

```
In [39]: # Load the matrix here
         user_item_matrix = pd.read_pickle('user_item_matrix.p')

In [40]: # quick look at the matrix
         user_item_matrix.head()
```

```
Out[40]: article_id  0.0  100.0  1000.0  1004.0  1006.0  1008.0  101.0  1014.0  1015.0  \
         user_id
         1           0.0   0.0    0.0     0.0     0.0     0.0    0.0    0.0     0.0
         2           0.0   0.0    0.0     0.0     0.0     0.0    0.0    0.0     0.0
         3           0.0   0.0    0.0     0.0     0.0     0.0    0.0    0.0     0.0
         4           0.0   0.0    0.0     0.0     0.0     0.0    0.0    0.0     0.0
         5           0.0   0.0    0.0     0.0     0.0     0.0    0.0    0.0     0.0

         article_id  1016.0  ...   977.0  98.0  981.0  984.0  985.0  986.0  990.0  \
         user_id             ...
         1              0.0   ...    0.0   0.0   1.0    0.0    0.0    0.0    0.0
         2              0.0   ...    0.0   0.0   0.0    0.0    0.0    0.0    0.0
         3              0.0   ...    1.0   0.0   0.0    0.0    0.0    0.0    0.0
         4              0.0   ...    0.0   0.0   0.0    0.0    0.0    0.0    0.0
         5              0.0   ...    0.0   0.0   0.0    0.0    0.0    0.0    0.0
```

```
article_id   993.0   996.0   997.0
user_id
1               0.0     0.0     0.0
2               0.0     0.0     0.0
3               0.0     0.0     0.0
4               0.0     0.0     0.0
5               0.0     0.0     0.0

[5 rows x 714 columns]
```

2. In this situation, you can use Singular Value Decomposition from numpy on the user-item matrix. Use the cell to perform SVD, and explain why this is different than in the lesson.

In [41]: *# Perform SVD on the User-Item Matrix Here*

       u, s, vt = np.linalg.svd(user_item_matrix)

In [42]: print(u.shape)
       print(s.shape)
       print(vt.shape)

```
(5149, 5149)
(714,)
(714, 714)
```

Provide your response here.

3. Now for the tricky part, how do we choose the number of latent features to use? Running the below cell, you can see that as the number of latent features increases, we obtain a lower error rate on making predictions for the 1 and 0 values in the user-item matrix. Run the cell below to get an idea of how the accuracy improves as we increase the number of latent features.

In [43]: num_latent_feats = np.arange(10,700+10,20)
       sum_errs = []

       for k in num_latent_feats:
           # restructure with k latent features
           s_new, u_new, vt_new = np.diag(s[:k]), u[:, :k], vt[:k, :]

           # take dot product
           user_item_est = np.around(np.dot(np.dot(u_new, s_new), vt_new))

           # compute error for each prediction to actual value
           diffs = np.subtract(user_item_matrix, user_item_est)

           # total errors and keep track of them
           err = np.sum(np.sum(np.abs(diffs)))
           sum_errs.append(err)

```
plt.plot(num_latent_feats, 1 - np.array(sum_errs)/df.shape[0]);
plt.xlabel('Number of Latent Features');
plt.ylabel('Accuracy');
plt.title('Accuracy vs. Number of Latent Features');
```


Accuracy vs. Number of Latent Features

4. From the above, we can't really be sure how many features to use, because simply having a better way to predict the 1's and 0's of the matrix doesn't exactly give us an indication of if we are able to make good recommendations. Instead, we might split our dataset into a training and test set of data, as shown in the cell below.

Use the code from question 3 to understand the impact on accuracy of the training and test sets of data with different numbers of latent features. Using the split below:

- How many users can we make predictions for in the test set?

- How many users are we not able to make predictions for because of the cold start problem?
- How many articles can we make predictions for in the test set?

- How many articles are we not able to make predictions for because of the cold start problem?

```
In [44]: df_train = df.head(40000)
         df_test = df.tail(5993)

         def create_test_and_train_user_item(df_train, df_test):
             '''
```

```python
            INPUT:
            df_train - training dataframe
            df_test - test dataframe

            OUTPUT:
            user_item_train - a user-item matrix of the training dataframe
                              (unique users for each row and unique articles for each column)
            user_item_test - a user-item matrix of the testing dataframe
                             (unique users for each row and unique articles for each column)
            test_idx - all of the test user ids
            test_arts - all of the test article ids

            '''
            user_item_train = create_user_item_matrix(df_train)
            user_item_test = create_user_item_matrix(df_test)

            test_idx = user_item_test.index
            test_arts = user_item_test.columns

            return user_item_train, user_item_test, test_idx, test_arts

        user_item_train, user_item_test, test_idx, test_arts = create_test_and_train_user_item(
```

In [51]: # Exploring
```python
         print(test_idx.shape)
         print(len(test_arts))
```

(682,)
574

In [46]: # Replace the values in the dictionary below
```python
         a = 662
         b = 574
         c = 20
         d = 0


         sol_4_dict = {
             'How many users can we make predictions for in the test set?': c,
             'How many users in the test set are we not able to make predictions for because of
             'How many articles can we make predictions for in the test set?': b,
             'How many articles in the test set are we not able to make predictions for because
         }

         t.sol_4_test(sol_4_dict)
```

Awesome job! That's right! All of the test articles are in the training data, but there are on

5. Now use the **user_item_train** dataset from above to find U, S, and V transpose using SVD. Then find the subset of rows in the **user_item_test** dataset that you can predict using this matrix decomposition with different numbers of latent features to see how many features makes sense to keep based on the accuracy on the test data. This will require combining what was done in questions 2 - 4.

Use the cells below to explore how well SVD works towards making predictions for recommendations on the test data.

```
In [52]: # fit SVD on the user_item_train matrix
         u_train, s_train, vt_train = np.linalg.svd(user_item_train)

In [56]: # decomposition to predict on test data
         print(u.shape)
         print(s.shape)
         print(vt.shape)


         num_latent_feats = np.arange(10,700+10,20)
         sum_errs_train = []
         sum_errs_test = []


         #Decomposition
         row_i = user_item_train.index.isin(test_idx)
         col_i = user_item_train.columns.isin(test_arts)

         u_test = u_train[row_i, :]
         vt_test = vt_train[:, col_i]

         users_predict = np.intersect1d(list(user_item_train.index),list(user_item_test.index))
         for k in num_latent_feats:
             # restructure with k latent features
             s_train_n, u_train_n, vt_train_n = np.diag(s_train[:k]), u_train[:, :k], vt_train[:
             u_test_n, vt_test_n = u_test[:, :k], vt_test[:k, :]

             # take dot product
             user_item_train_preds = np.around(np.dot(np.dot(u_train_n, s_train_n), vt_train_n))
             user_item_test_preds = np.around(np.dot(np.dot(u_test_n, s_train_n), vt_test_n))

             # compute error for each prediction to actual value
             diffs_train = np.subtract(user_item_train, user_item_train_preds)
             diffs_test = np.subtract(user_item_test.loc[users_predict,:], user_item_test_preds)

             # total errors
             err_train = np.sum(np.sum(np.abs(diffs_train)))
             err_test = np.sum(np.sum(np.abs(diffs_test)))

             sum_errs_train.append(err_train)
             sum_errs_test.append(err_test)
```

22

```
(5149, 5149)
(714,)
(714, 714)
```

```
In [61]: f, ax = plt.subplots(figsize=(16, 8))
         ax.set_ylabel('Test Accuracy')
         ax.plot(num_latent_feats, 1 - np.array(sum_errs_test)/df.shape[0])
         ax.tick_params(axis='y')
         ax.set_title("Accuracy vs # of latent features")
```

```
Out[61]: Text(0.5,1,'Accuracy vs # of latent features')
```

Accuracy vs # of latent features

6. Use the cell below to comment on the results you found in the previous question. Given the circumstances of your results, discuss what you might do to determine if the recommendations you make with any of the above recommendation systems are an improvement to how users currently find articles?

Your response here.

We see that the accuracy for the test data decreases with an increase in latent features. This is mostly due to over fitting of the data. It is best to keep the latent features relatively low.

Conduct A/B test to solve the cold start problem and evaluate recommendation engine performance. Recommend articles to one group using the engine and to the other with most popular articles. Compare click-through rates (maybe using cookies) to measure increase in clicks. If a statistically significant(p-value test) rise in clicks is observed, the recommendation engine is deemed successful.

Extras Using your workbook, you could now save your recommendations for each user, develop a class to make new predictions and update your results, and make a flask app to deploy your results. These tasks are beyond what is required for this project. However, from what you learned in the lessons, you certainly capable of taking these tasks on to improve upon your work here!

1.2 Conclusion

Congratulations! You have reached the end of the Recommendations with IBM project!

Tip: Once you are satisfied with your work here, check over your report to make sure that it is satisfies all the areas of the rubric. You should also probably remove all of the "Tips" like this one so that the presentation is as polished as possible.

1.3 Directions to Submit

Before you submit your project, you need to create a .html or .pdf version of this notebook in the workspace here. To do that, run the code cell below. If it worked correctly, you should get a return code of 0, and you should see the generated .html file in the workspace directory (click on the orange Jupyter icon in the upper left).

Alternatively, you can download this report as .html via the **File** > **Download as** submenu, and then manually upload it into the workspace directory by clicking on the orange Jupyter icon in the upper left, then using the Upload button.

Once you've done this, you can submit your project by clicking on the "Submit Project" button in the lower right here. This will create and submit a zip file with this .ipynb doc and the .html or .pdf version you created. Congratulations!

```python
In [62]: from subprocess import call
         call(['python', '-m', 'nbconvert', 'Recommendations_with_IBM.ipynb'])

Out[62]: 0

In [ ]:
```